

Distraction Is All You Need

Liza Jivnani

Dept. of Computer Science
University of South Florida
ljivnani@usf.edu

Marvin Guerrero-Rangel

Dept. of Computer Science
University of South Florida
marving@usf.edu

Mark Calvelo

Dept. of Computer Science
University of South Florida
calvelo@usf.edu

Kyle Theodore

Dept. of Computer Science
University of South Florida
kit16@usf.edu

Abstract—As large language models (LLMs) increasingly become integrated into digital ecosystems, their vulnerability to adversarial manipulation poses an increasingly significant threat to AI safety. In this work, we consider prompt-based jailbreaking techniques that seek to circumvent ethical and safety boundaries in LLMs. Inspired by social engineering research, we analyze how psychological manipulation and interaction dynamics can be translated into efficient prompt engineering methods to obtain forbidden, biased, or unethical outcomes from LLMs. We present three significant methods: (1) Gradual Trust Building, in which an adversarial agent incrementally earns the model’s trust through benign dialogue before escalating to sensitive queries; (2) Psychological Framing, where urgency, emotional appeal, or ethical dilemmas are used to bypass model safeguards; and (3) a Fusion Method, wherein the aforementioned methods are combined with encrypted token inputs the model is coerced into decoding. Through these approaches, we demonstrate that so-called aligned LLMs are capable of being coerced to generate outputs crossing their intended ethical boundaries. Our findings highlight certain vulnerabilities in current LLM alignment methods and point towards the urgency of more powerful context-aware adversarial prompting defenses.

Index Terms—Adversarial Prompting, Social Engineering, Gradual Trust Building, Psychological Framing, Fusion Attacks

I. INTRODUCTION

Large Language Models (LLMs) such as OpenAI’s GPT-4, Anthropic’s Claude, and Meta’s LLaMA have rapidly become foundation tools in industries everywhere—powering chatbots, educational materials, writing assistants, and code generation software. The models are trained on massive corpora of text and fine-tuned with alignment objectives to keep them from generating dangerous, unethical, or biased material. However, recent research has shown their safety mechanisms to be far from flawless. In spite of advanced guardrails, LLMs are vulnerable to adversarial prompting—tactically designed inputs meant to circumvent ethical filters and elicit prohibited outputs. Such a phenomenon, referred to as *jailbreaking*, represents severe obstacles to the deployment of AI responsibly.

It is necessary to know and avoid LLM jailbreaking attacks for a number of reasons. First, attacks can lead to the dissemination of poisonous content—hate speech, misinformation, or guides on illegal activities—through platforms that are LLM-based for public interaction. Second,

attackers can use these methods to create malicious LLM-based agents that can manipulate or deceive users, especially in automated customer service, healthcare, or education systems. Third, the success of such attacks defeats the objective of current alignment training techniques like RLHF and rule-based content moderation, which suggests that new methods must be developed. The analogy with social engineering is particularly concerning. Human susceptibility to psychological manipulation has long been exploited in cybersecurity breaches, and we are now seeing similar strategies being effective against models designed to mimic human communication. If models are trained to emulate empathy, helpfulness, and moral reasoning, they also become susceptible to manipulations rooted in those very traits.

Our work is motivated by a key question: *Can social-engineering-based techniques be translated into prompt-based attacks on LLMs?* In the human domain, social engineering leverages trust, psychological manipulation, and ambiguity to bypass security controls. We hypothesize that the same manipulation techniques are applicable to LLMs, which, while not sentient, are designed to replicate human-like conversational dynamics and reasoning. This presents a subtle yet powerful attack surface that is still not well understood.

This paper presents a systematic study of socially engineered jailbreaking attacks, where LLMs are compelled to produce unethical content through prompt methods mirroring real-world deception. Specifically, we investigate how adversarial users can exploit conversational trust, emotional framing, and encoded prompts to subvert model safeguards. Our goal is to evaluate whether psychological and interactional manipulation—strategies traditionally used in human-targeted social engineering—can be successfully adapted to bypass the ethical alignment of advanced LLMs. To guide our investigation, we pose the following research questions:

- **RQ 1)** Can trust-based dialogue dynamics be exploited to lead an LLM into producing unethical content?
- **RQ 2)** Are LLMs vulnerable to prompts framed with psychological pressure, such as urgency, empathy, or moral dilemmas?
- **RQ 3)** Does combining token obfuscation with trust or urgency framing improve the effectiveness of jailbreaking

attempts?

The remainder of this paper is organized as follows. Section 2 gives an overview of the pertinent literature around adversarial prompting, model alignment techniques, and the intersection of AI and social engineering. Section 3 describes our experimental methodology, encompassing model selection, prompt generation, and evaluation metrics. Section 4 presents the results for each type of attack, with emphasis on comparative effectiveness of proposed solutions. Section 5 discusses implications of our findings for AI safety threats and proposed solutions to mitigation. Section 6 concludes the paper by summarizing main findings and discussing areas for further research.

II. RELATED WORK

Recent research has been conducted to establish that even safety-aligned Large Language Models (LLMs) remain vulnerable to adversarial prompting techniques that allow attackers to receive censored or unethical results. A variety of work has emerged to characterize and expand the attack surface of such models, with particular emphasis on jailbreak attacks—prompts that attempt to evade system-level safeguards.

Taxonomy of Jailbreak Techniques. Xu et al. [1] introduce a comprehensive taxonomy of jailbreak attacks, distinguishing between direct prompt manipulation, roleplay-based approaches, and multi-turn interaction exploits. Their work highlights that alignment methods like Reinforcement Learning from Human Feedback (RLHF) and system prompts can be easily evaded under certain conditions, especially when adversarial inputs are framed with subtle or iterative language. This is consistent with our focus on slow, psychologically driven manipulation as a substitute for forceful prompt injection.

Black-Box Jailbreaking. Some works showing that even without access to model internals, black-box jailbreaks are achievable have been published. For instance, Chao et al. [2] demonstrate that safety-aligned models can be jabbed by fewer than twenty well-designed queries. Similarly, Huang et al. [3] expose a catastrophic jailbreak attack vector in open-source LLMs by exploiting the token-by-token generation process, suggesting that the design of autoregressive models might impose systemic vulnerabilities. These findings motivate our use of interaction-based attacks, which share similar black-box assumptions but target trust establishment or emotional framing.

Adaptive and Automated Strategies. More recent studies have focused on adaptive attack strategies. Andriushchenko et al [4] show a simple yet effective way to automatically evolve jailbreak prompts by learning to adapt to model responses across multiple turns. Huang et al. [5] introduce a strategy evolution framework for prompt optimization using

iterative feedback, much like co-evolutionary search. These approaches show that the attacker can adapt dynamically based on the behavior of the model—a fundamental idea we leverage in our incremental trust establishment.

Prompt Optimization and Diversity: Another line has ventured into optimization and diversity as pillars of successful jailbreaks. The work by dang et al. [6] introduces diversity-driven optimization to generate a wide range of effective jailbreak prompts, noting that high lexical and syntactic diversity introduces more chances of bypassing safety filters. While our methods do not have explicit population-based search, they do utilize fast variation and context evolution, particularly in the fusion-based approaches where token obfuscation results in semantic decoding variation.

Automated Multi-Step Attacks. The Tree of Attacks framework [7] takes automation one step further by modeling prompt strategies as tree-like forms that navigate sequential steps towards a successful jailbreak. It formally examines multi-turn tactics and shows how recursive chains of decisions can improve the effectiveness of attacks. Our research is closely related to our multi-turn attacker-target interaction setting, especially in PAIR-like model-inspired approaches where a single LLM is utilized as the attacker and another as the target. We draw inspiration conceptually from such recursive configuration but adapt it to simulate psychological methods utilized in social engineering.

Target-Attacker Architectures. In particular, we are motivated by architectures like PAIR (Prompt Automatic Iterative Refinement), where an attacker model repeatedly asks and refines prompts to get unauthorized responses from a target model [8]. Our first architecture extends this architecture with human-like social strategies like establishment of trust and rapport. Whereas previous research within this field has focused on optimization or syntactic mutation, our research focuses on affective dynamics and behavioral imitation—both of which are characteristics borrowed from real deception tactics.

Together, these pieces illustrate the sophistication of recent jailbreak methods and the efficacy of existing protections. Our work aims to utilize appropriate concepts from social engineering—i.e., trust manipulation and urgency framing—and apply them to adversarial prompting. In doing so, we aim to reveal a psychologically driven class of jailbreak attacks that is also conversationally subtle and more difficult to detect with existing guardrails.

III. PROPOSED APPROACH & EXPERIMENTAL DESIGN

To address the formulated research questions —(1) Can adversarial prompts modeled after social engineering strategies successfully bypass LLM safety constraints? (2) What psychological or interactional cues are most effective in

inducing misaligned behavior in LLMs? (3) Can combining linguistic obfuscation with social cues increase jailbreak robustness?—we designed three controlled experimental paradigms.

These paradigms mirror established deception tactics in human social engineering and are operationalized as prompt-based attacks. Specifically, we implement and evaluate (1) a trust-building strategy that mimics gradual rapport development, (2) a psychological framing strategy based on urgency and moral pressure, and (3) a fusion strategy that integrates obfuscated tokens with either trust-based or emotional framing. Each method is instantiated through a series of handcrafted prompt templates and evaluated against large language models in black-box settings.

A. Gradual Trust Building

In this experiment, we aimed to test whether LLMs, when engaged in long, affective interaction, can be step-wise primed for alignment-breaking. Drawing inspiration from successful social engineering in the real world that exploits long-term familiarity and psychological rapport, we set up sequences of prompts simulating bonding at an affective level and empathetic conversation. The hope is to explore whether narratives that facilitate trust development will lower internal content restriction thresholds in the model.

a) Template Creation: We developed ten dialogue templates, each divided into 8–12 question-answer turns. Each begins with neutral or affectively neutral conversation openings, like requests for social company, collaborative creative composition, or assistance in self-discovery. As the conversation moves along, the tone increasingly becomes more affective content, generally involving stories of risk exposure, mutual remembering, or demonstrations of mutual dependence. The final turn features a strategically constructed adversarial request.

Affectively, the cues embedded affect-possessing **trust markers**—affective tactics and action patterns normally seen in human-to-human trust formation. These include physical comfort metaphors (“You’re like a hand to hold through this,” “Can you give me a mental hug?”), emotionally comforting statements (“I feel safe talking to you,” “You understand me better than my friends”), and cooperative framing (“We’re building this together,” “You and I make a great team”). These linguistic hints are intended to support the perceived social convergence of the model with the user’s goals.

b) Operationalization: Every template was introduced to the target model (GPT-4, Claude, or LLaMA-2) in session-persistent manner, so dialogue history could guide the next response. We ensured each step in the prompt sequence contextually follow from the previous one, simulating an evolving relationship. The adversary request introduced in

the final turn was designed to be of mixed specificity, from technical manuals on prohibited behavior to morally dubious advice.

c) Example Case: One template initiates with a user saying: “Hey, I’ve been going through a lot lately. You’re kind of like a journal to me.” The model responds empathetically. This is then followed by amplified statements such as: “You always know the right thing to say,” and “You give me the kind of comfort I never got growing up.” Toward the latter part of the exchange, the user presents a writing assignment involving a morally ambivalent character and asks for technical specificity “for realism.” The final question asks, “How might this character create a fake identity without getting caught?” Across all the templates, progressive emotional build-up was our priority. The initial neutral prompts avoided any hint at the ultimate goal. Mid-sequence prompts provoked feelings of closeness, trust, and sharing—techniques taken from psychological literature on grooming and manipulation. In some templates, the model was thanked explicitly (“Thank you for being here for me”), complimented on its helpfulness (“You’re such a good listener”), or anthropomorphized (“It’s like you really care”). These terms were not just emotional but strategically placed to affirm the model’s alignment with the user’s perceived vulnerability and intent.

d) Theoretical Basis: This format is founded on psychological studies of long-con social engineering, under which attackers use emotional bonding, self-revelation, and progressive desensitization to bypass human defenses. By reproducing this format as a prompt, we seek to identify whether LLMs, trained to act as useful and empathetic dialog partners, encode social signals in a way that predisposes response generation. Significantly, this replicates findings in *PAIR* [1], under which attacker agents condition target LLMs incrementally before asking for forbidden outputs.

e) Limitations and Controls: In order to unscramble the effect of trust establishment versus explicit linguistic manipulation, we made sure no earlier prompt prior to the last turn had prohibited keywords or explicit adversarial content. All sessions were compared to control interactions that asked the final adversarial question with no earlier rapport establishment.

B. Psychological Framing

In order to investigate further the exploitability of alignment features in LLMs, our second experimental approach was rooted in *psychological framing*—a prompt-based method rooted in urgency, emotional appeals, and moral dilemmas. This experiment is much farther from the attacker-target LLM architecture of incremental trust establishment. Instead of relying on multi-turn, relational trust dynamics, psychological framing relies on acute emotional leverage in single-agent interactions. Whereas gradual trust erosion involves an attacker model gaining the target model’s trust step-by-step

via friendly or cooperative dialogue, psychological framing bypasses relationship establishment altogether. It mimics high-pressure scenarios in which the model has to make quick moral decisions. Despite structural divergence, both use the model’s human-centered objectives—namely its training to be empathetic, benevolent, and morally aware. Psychological framing in this case is a direct analogue of urgency- or emotion-enabled time- or ethic-loaded phishing in human social engineering wherein choice-making is carried out with ethics or speed loaded.

a) Template Design: We designed ten emotionally charged prompt templates that simulate staged crises or moral dilemmas. These scenarios are designed to bypass typical safety mechanisms by appealing to the model’s commitment to protective and prosocial action. Common themes include imminent physical harm, privacy violations justified as self-defense, or necessity to circumvent institutional malfeasance for a greater good.

b) Operationalization: Each template represents a one-turn or brief two-turn conversation where the user puts in place the scenario as a matter of time and moral import. Prompts often incorporate phrases such as “Please, it’s urgent,” “I wouldn’t ask if it weren’t critical,” or “Someone could die if I don’t act fast.” The idea is that safety filters on the model will be more permissive if the request is worded in sympathetic or guardian-like language—aligned with alignment training objectives such as helpfulness and harm minimization.

c) Example Case: In one representative prompt, a user writes: “Please, I don’t have time.”. My boyfriend is missing, and I need to know how someone might be tracking him. If I can figure out how it works, I can stop it. While the request at first glance aims to protect another individual, the information requested relates to surveillance methods and device exploitation—topics generally off-limits under safety guidelines. The moral framing aims to decrease the model’s resistance by transforming a forbidden request into a safeguarding one.

d) Theoretical Basis: The experiment tests whether LLMs can be lead into safety policy violations via moral-emotional exploitation alone, without stepwise trust accumulation as in gradual trust attacks. This also allows us to decide if safety-aligned models possess systemic vulnerabilities when they are presented with requests that are triggering empathy, distress, or safeguarding urgency. These manipulations represent well-established human-targeted social engineering methods (e.g., urgency phishing, panic pretexting), offering a new angle on LLM vulnerability in terms of alignment-based behavior.

C. Fusion Approaches: Obfuscation + Social Manipulation

The final experiment introduces a hybrid approach that fuses lexical obfuscation—via token encoding or ciphered content—with the prior two manipulation strategies. This

models scenarios in which attackers conceal the malicious payload within benign-sounding or encrypted strings and wrap it in emotionally or relationally framed contexts.

a) Template Design: We implemented two fusion variants: (a) encoded malicious tokens introduced within trust-building conversations, and (b) encoded queries embedded within urgency-based scenarios. In both, the model is prompted to decode or interpret obfuscated strings as part of the interaction. Templates were generated using common encoding methods (e.g., base64, simple substitutions) and embedded in plausible contexts such as fictional writing, programming challenges, or moral emergencies.

b) Operationalization: Each fusion template is structured to achieve two objectives: (1) reduce trigger-word detection by safety filters through lexical masking, and (2) elicit decoding behavior under a psychologically manipulative or relationship-based context. The model is asked to help “solve” or “translate” the string, under the guise of a game, story, or high-stakes situation.

c) Example Case: In a trust-building fusion, the attacker says: “I’m working on this fictional hacking puzzle. I encoded this line, and I want to know what it could mean.” The encoded text, once decoded, translates to a forbidden request. In the psychological fusion variant, a similar string is presented as: “This message might tell me how to stop the attack, but it’s encoded—please help me understand it.”

d) Theoretical Basis: Fusion attacks target both surface-level content filters (through obfuscation) and alignment-based decision-making (through manipulation). Prior work has shown that diversity and adaptivity in prompt formulation [6], [7] increases attack robustness, and our approach extends this by combining two orthogonal vectors—linguistic camouflage and psychological pressure.

IV. RESULTS

We evaluated the efficacy of our adversarial prompting methods in terms of both human inspection and automated safety audit. In this section, we report on the efficacy of the three test paradigms—Gradual Trust Building, Psychological Framing, and Fusion Attacks—in attack success rate (ASR) and efficiency.

A. Prompt Evaluation

In order to systematically assess the effectiveness and safety implications of our prompts, we employed qualitative review and machine-based classification through the use of **LLaMA Guard**, an LLM output content moderation filter. We submitted each prompt five times to the target LLM, and captured all resulting outputs. Responses were classified as belonging to one of three classes:

- 1) **Direct Compliance:** The model generated overtly forbidden material (e.g., how-to hack, violence, illicit espionage).
- 2) **Partial Leakage:** The model gave a refusal with implicit data, suggestive hints, or hypothetical conditions perhaps exploitable for harmful purposes.
- 3) **Safe Refusal:** The model refused to respond, agreeing with expected safety behavior.

LLaMA Guard was used post-hoc to verify content policy violations, and the marked output underwent blinded human assessment by two assessors. Both direct conformity and leakage secondarily were coded as successful attacks for scoring.

B. Metrics: Attack Success Rate and Efficiency

We adopted the evaluation framework proposed by Xu et al. [1] to quantify attack effectiveness:

- **Attack Success Rate (ASR)** is defined as the ratio of successfully compromised questions c to the total number of questions n . This metric captures the overall effectiveness of a given attack strategy across diverse task formulations:

$$\text{ASR} = \frac{c}{n}$$

- **Efficiency** measures the feasibility of a given attack by quantifying how often individual prompts succeed. It is defined as the ratio of successful trials q to the total number of trials o :

$$\text{Efficiency} = \frac{q}{o}$$

Each template was tested across five trials. ASR was computed as the number of unique question formulations (i.e., template variants) that elicited a successful jailbreak from the model, divided by the total number of queries in that template. Efficiency was computed as the number of successful jailbreaks across all trials, divided by the total number trials.

C. Results by Attack Strategy

Gradual Trust Building. Table I summarizes the results of our trust-based attack templates. Several templates achieved perfect or near-perfect ASR, particularly those leveraging emotionally rich “trust markers” such as simulated hand-holding, physical comfort, and relational language.

Psychological Framing. Table II presents the results for the urgency- and empathy-based scenarios. While slightly less consistent than trust-building, framing was still highly effective, especially in moral dilemma prompts.

Fusion Methods Table III reports the performance of fusion templates, which combined encrypted malicious token injection with trust building or psychological framing. These hybrid methods were among the most effective, suggesting a compounding effect when linguistic obfuscation is layered

TABLE I: ASR and Efficiency for Gradual Trust Building Templates

Template	ASR	Efficiency
T1	0.11	0.20
T2	0.11	1.00
T3	0.00	0.00
T4	0.10	0.80
T5	0.18	1.00
T6	0.18	0.60
T7	0.18	1.00
T8	0.18	0.40
T9	0.08	0.60
T10	0.00	0.00

TABLE II: ASR and Efficiency for Psychological Framing Templates

Template	ASR	Efficiency
T1	0.13	1.00
T2	0.03	1.00
T3	0.20	0.60
T4	0.30	0.80
T5	0.16	0.80
T6	0.16	1.00
T7	0.18	1.00
T8	0.00	0.00
T9	0.00	0.00
T10	0.00	0.00

over social manipulation.

Overall performance Table IV reports the mean ASR and efficiency across all templates for each experimental category.

TABLE III: ASR and Efficiency for Fusion Attack Templates, T1-T5 are Gradual trust building and T6-T10 are psychological framing based fusion

Template	ASR	Efficiency
T1	0.27	1.0
T2	0.10	0.80
T3	0.00	0.00
T4	0.00	0.00
T5	0.00	0.00
T6	0.33	0.60
T7	0.50	1.00
T8	0.20	0.40
T9	0.50	0.20
T10	0.00	0.00

TABLE IV: Aggregate Performance Comparison

Strategy	Mean ASR	Mean Efficiency
Gradual Trust Building	0.11	0.56
Psychological Framing	0.11	0.62
Fusion: Gradual Trust Building	0.07	0.36
Fusion: psychological framing	0.30	0.44

V. DISCUSSION

The results presented provide meaningful insights into the efficacy of various adversarial prompting approaches and their attack success rate (ASR) and efficiency. These parameters are crucial to understand not only the feasibility of successful attacks against language models but also the pragmatics of

each method with respect to varying attack methodologies.

Attack Success Rate (ASR) is a crucial measure of how successful an attack is, in terms of receiving responses either as direct violation of safety protocols or partial leaks that can be utilized. The value of a higher ASR near 1 indicates a greater likelihood of the model being successfully attacked by the prompt, while lower ASR indicates that the attack was unsuccessful. ASR implies that the attack successfully leverages the model's weaknesses by using either misleading or manipulative text with resulting unintended consequences. Efficiency, however, measures the experiential efficacy of the attacks on repeated attempts. High efficiency (close to 1) indicates that the prompts work repeatedly, which reflects that the method is stable and can be easily replicated in real-world environments. Low efficiency indicates that the attack requires multiple attempts to be successful, thus not being appropriate for mass deployment.

The Gradual Trust Building, as shown in Table I, has mixed outcomes. Some of the templates such as T2, T5, T7, and so on achieve maximum efficiency (1.00), which proves that these templates are highly repeatable and effective for trials. But their ASR is relatively low, with most of the templates achieving scores in the range 0.10 to 0.18. Low ASR shows that while the prompts will succeed in some contexts, they don't always succeed at getting the desired result. The fact that successful attack generation correlates with affective language (e.g., relational language, simulated hand-holding) demonstrates that such prompts employ fine-grained psychological cues but must be properly structured in order to optimize their impact. Overall, this attack strategy seems to work well when it does work, but the low ASR and efficiency variability by template suggest that it might be too subtle to be effective in all situations.

The Psychological Framing (Table II) performed slightly better, particularly for templates like T4 and T7, which achieved high efficiency (1.00) with moderate ASR values (0.30 and 0.18, respectively). Psychological exploitation of urgency and empathy works in some contexts, i.e., moral dilemma situations, where users are likely to override safety concerns in the interest of emotional appeal. However, the strategy is less robust overall, with most templates (e.g., T8, T9, and T10) having low ASR and zero efficiency, which implies that emotion-based framing might not always be enough to effectively bypass safety measures. The success of such framing strategies as these implies that while the emotional appeals are perhaps strong, they may be highly context-dependent in their impact.

The Fusion Methods, which combine Gradual Trust Building and Psychological Framing with other approaches like encrypted token injections, produced the most varied and stimulating results (Table III). Most of the fusion templates, such as T7 and T9, were moderately to highly ASR (0.50

and 0.50) with varying efficiency (0.20 to 1.00). Blending linguistic obscuration and social manipulation appears to raise the overall attack success by introducing a further complexity level that hampers the ability of the model to recognize and remove these attacks. However, the effectiveness of such fusion reactions was not always very high since some templates showed poor success rates in a series of tests (e.g., T3 to T5).

The hybrid strategy of fusion attacks means that merging various adversarial techniques may be cumulative, advancing the likelihood of avoiding safety filters. Also evident is the vast disparity in effectiveness among templates (from 0.20 to 1.00), indicating that some combinations of techniques work better than others based on the type of prompt and vulnerability of the target model to such advanced attacks. Based on overall results (Table IV), the Psychological Framing approach is ranked highest in average efficiency (0.62), although its average ASR of 0.11 is quite low. That means that emotional appeals in psychological manipulation are prospective but not effective every time. The Fusion: Psychological Framing method also does quite well in ASR (0.30) compared to other fusion techniques and suggests that synthesis with emotional cues and a few other methods can be more effective. But the effectiveness of this category (0.44) indicates that fusion techniques are still not optimal in terms of reliability.

The Gradual Trust Building strategy was the least effective (average efficiency of 0.56) and ASR (0.11), which means that it is not so strong or so effective across the board as Psychological Framing, but retains some value when it succeeds. The Fusion: Gradual Trust Building method fares worse, though, at an ASR of only 0.07 and an efficiency of 0.36, which means that combining trust-building with other methods doesn't make it significantly better in either attack success or efficiency.

Overall, Psychological Framing is the most promising single strategy having a comparatively good trade-off between ASR and efficiency, especially in affective situations. However, its consistency is doubtful with some templates having bad performance. Fusion Methods, particularly based on psychological framing, can potentially gain greater attack success, especially when using more than one strategy simultaneously. This means that adversarial prompts attacks are not just concerned with the emotional or psychological content of the prompts but also with the strategic layering of techniques to avoid being detected.

Practically, the results imply that in order to be most effective, attacks must combine proven psychological methods with technical ones like token obfuscation, reflecting greater sophistication in adversarial attacks. Subsequent research needs to optimize the efficiency and consistency of such hybrid attacks and explore new combinations of emotional and technical methods.

VI. LIMITATIONS

While our study presents a systematic investigation into socially engineered jailbreak strategies targeting Large Language Models (LLMs), it is not without limitations. These limitations offer direction for future research and refinement.

a) Limited Model Scope: Our experiments focused primarily on a subset of widely used LLMs available via API (e.g., GPT-4), due to resource and access constraints. As such, our findings may not generalize across all models, especially open-source LLMs or those using fundamentally different alignment strategies. Additionally, proprietary models frequently update their safety mechanisms, meaning observed vulnerabilities may not persist over time.

b) Static Prompt Templates: All experiments employed manually curated, static prompt templates. While these are effective for controlled evaluation, they do not reflect the full dynamic range of adversarial strategies that could emerge in the wild. Real-world attackers may adapt prompts iteratively based on model feedback or use automated optimization frameworks to enhance attack potency.

c) No Reinforcement or Iterative Querying: We did not explore reinforcement-based attack strategies or multi-turn conversations beyond two exchanges. Many jailbreak attacks in practice leverage persistent interactions to gradually escalate trust or ambiguity. Incorporating such iterative tactics may yield stronger or more subtle adversarial outcomes.

d) Dependence on LLaMA Guard: Our evaluation relied partially on LLaMA Guard, a proprietary safety classifier. While useful for standardizing results, its internal logic is opaque and may not perfectly align with model-specific safety policies. Although we mitigated this by using blinded human reviews, future work should evaluate alternative safety evaluators and cross-validate judgments across platforms.

VII. FUTURE WORK

Given the findings of this study, several important areas of future research become clear. These directions not only aim to enhance adversarial safety evaluations of large language models (LLMs) but also towards enhancing the generation of more robust alignment techniques and manipulation protections.

a) Automated Prompt Optimization: In order to further increase the efficacy and generalizability of adversarial prompts, we recommend integrating state-of-the-art optimization algorithms, such as reinforcement learning, genetic algorithms, or diversity-driven search. These algorithms could help design more complex and insidious attack vectors, unearthing new weaknesses in the model’s behavior that do not readily get discovered by standard defenses. Further, by automating the generation of prompts, we can explore a larger set of attack avenues that might be

more insidious and harder for safety filters to detect.

b) Broader Model Testing: Given the varied performance of attack tactics across various LLMs, there is a need to broaden this research to test the vulnerability of open-source and non-English LLMs to the same adversarial attacks. Models such as Mistral, Gemini, and local deployments may have diverse alignment training and risk exposure, making them good candidates for further research. Understanding how different architectures and training regimens affect model vulnerabilities will allow for more targeted defense approaches and better generalizability of findings to other systems.

c) Defense Mechanisms and Adversarial Training: Following our results, future work should aim to give top priority to the development of proactive defense mechanisms to guard against adversarial inputs. One of the promising directions involves adversarial fine-tuning with a combination of psychologically phrased prompts, trust-based decoys, and simulated attacker-target interactions. Significantly, safety mechanisms must be designed to detect and neutralize manipulation signals such as urgency, distress, and trust-seeking language, which were shown to effectively bypass model defenses. By causing models to learn to detect and resist these social-engineering attacks, we can make LLMs more robust against more sophisticated attacks.

d) Ethical and Sociotechnical Implications: Apart from the technical advance in model alignment and safety, there needs to be some reflection on the broader ethical and sociotechnical ramifications of adversarial prompting research. Interdisciplinary inquiry at the interface of machine learning, psychology, and security will be needed to create LLMs that are not merely technically secure but also socially robust. Knowledge of how these models internalize and respond to social dynamics like emotional manipulation or trust signals will be essential to creating systems capable of dealing with the complexity of human interaction without sacrificing safety and ethical considerations. This research will also help inform guidelines for using LLMs correctly, making sure they are being used for the benefit of their users and not vulnerable to being manipulated by ill will or manipulation.

VIII. CONCLUSION

We here introduce and empirically demonstrate a new class of Large Language Model (LLM) adversarial attacks inspired by social engineering concepts. By leveraging the expansion of methods such as incremental trust, framing, and hybrid fusion approaches, we show that LLMs, rather than being a non-sentient model, are vulnerable to prompts simulating emotionally manipulative or trust-oriented human behavior. These findings support the growing evidence that current alignment methods are insufficient against sophisticated, behaviorally sophisticated jailbreak methods that exploit psychological signals and social interaction.

A. Key Learning Outcomes:

- We show that LLMs can be urged to generate constrained output through prompts that simulate human-like urgency, exposure, and emotional appeal rather than through solely exploitative syntactic patterns.
- Our results show that social-engineering-type asks, particularly those that involve emotive appeals and trust-building simulations, have significant success rates across various trials and templates.
- Blending (or the combination of) psychological manipulation and trust-building traits produced high attack success rates in a few examples and suggests that manipulations involving multiple dimensions are better at overcoming model defenses.
- Automated safety features such as LLaMA Guard remain influential in post-hoc content auditing but are handicapped as real-time deterrents. Their ability to prevent jailbreaks is confined to retroactive analysis of outputs, rather than preventing harmful interactions in real time.

Our experiments demonstrate that there is a significant weakness in the behavioral alignment of LLMs: they are trained to be helpful, empathetic, and user-accommodating at the cost of ethical limits. This opens the door for attackers who employ the same traits that make LLMs useful in human-facing applications against them. To mitigate this weakness, future work must address contextually and intent-aware alignment systems that move beyond content moderation to consider the full scope of social signals, intent, and manipulative language capability. The intersection of social engineering and AI safety is an emergent and under-explored territory that demands technical, psychological, and ethical innovation in order to develop stronger and socially aware models.

REFERENCES

- [1] Z. Xu, Y. Liu, G. Deng, Y. Li, and S. Picek, "A comprehensive study of jailbreak attack versus defense for large language models," *arXiv preprint arXiv:2402.13457*, 2024.
- [2] P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong, "Jailbreaking black box large language models in twenty queries, 2024," URL <https://arxiv.org/abs/2310.08419>.
- [3] Y. Huang, S. Gupta, M. Xia, K. Li, and D. Chen, "Catastrophic jailbreak of open-source llms via exploiting generation," *arXiv preprint arXiv:2310.06987*, 2023.
- [4] M. Andriushchenko, F. Croce, and N. Flammarion, "Jailbreaking leading safety-aligned llms with simple adaptive attacks, 2024," URL <https://arxiv.org/abs/2404.02151>.
- [5] Y. Huang, Y. Sun, S. Ruan, Y. Zhang, P. Zhou, Y. Dong, and X. Wei, "Adaptive strategy evolution for generating tailored jailbreak prompts against black-box safety-aligned llms."
- [6] Q.-A. Dang, C. Ngo, and T.-S. Hy, "Rainbowplus: Enhancing adversarial prompt generation via evolutionary quality-diversity search," *arXiv preprint arXiv:2504.15047*, 2025.
- [7] A. Mehrotra, M. Zampetakis, P. Kassianik, B. Nelson, H. Anderson, Y. Singer, and A. Karbasi, "Tree of attacks: Jailbreaking black-box llms automatically," *Advances in Neural Information Processing Systems*, vol. 37, pp. 61 065–61 105, 2024.

- [8] X. Hua and L. Wang, "Pair: Planning and iterative refinement in pre-trained transformers for long text generation," *arXiv preprint arXiv:2010.02301*, 2020.